Data augmentation for 10 s response acceleration

(a) Training-only cutout on the raw response

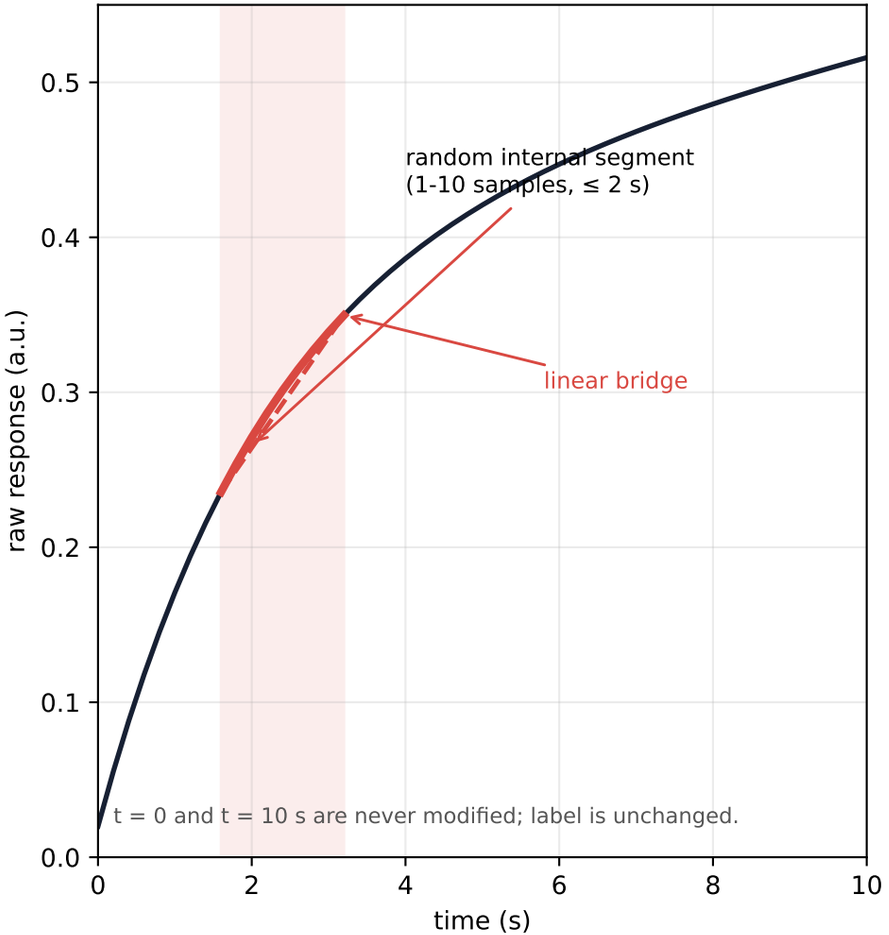

(b) Multi-prefix supervision pipeline

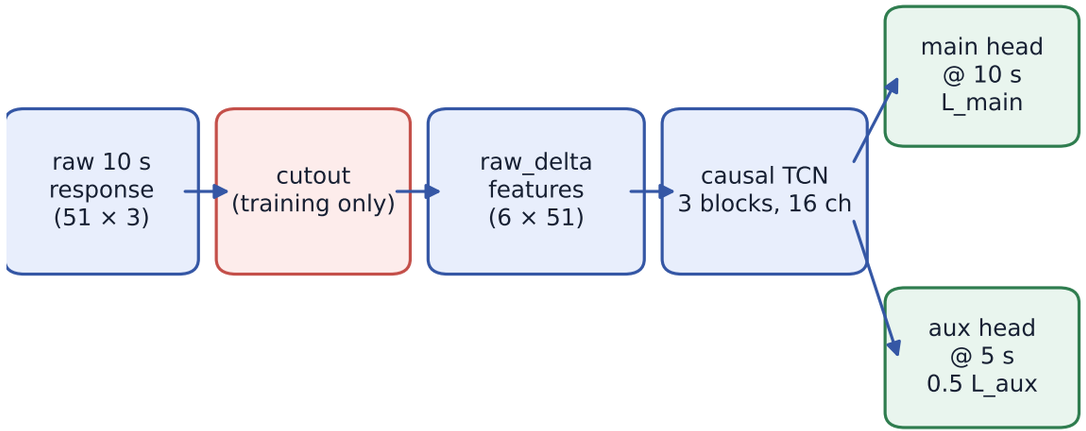

Validation / test / inference

clean 10 s response → TCN → main head → 3-seed average

No cutout, no auxiliary loss, no future samples.

total training loss

L_main + 0.5 L_aux